

An Introduction to Concurrency in Python

Natalia Maniakowska

skygate

16-01-2020

Outline

- 1 Introduction
 - Definitions
 - Different Types of Concurrency – More Details
- 2 Concurrency in Python
- 3 GIL

What is concurrency, actually?

- Medieval Latin *concurrentia* "a running together"
(cf. English *current*, *concur*)
- so: simultaneous occurrence...
- ...but is it **really** simultaneous? What does it mean, anyway?

What is concurrency, actually?

- Medieval Latin *concurrentia* "a running together"
(cf. English *current*, *concur*)
- so: simultaneous occurrence. . .
- . . . but is it **really** simultaneous? What does it mean, anyway?

What is concurrency, actually?

- Medieval Latin *concurrentia* "a running together"
(cf. English *current*, *concur*)
- so: simultaneous occurrence. . .
- . . . but is it **really** simultaneous? What does it mean, anyway?

Concurrency vs. Parallelism

https://wiki.haskell.org/Parallelism_vs._Concurrency

Disclaimer: Not all programmers agree on the meaning!

Definition

A **parallel program** is one that uses a multiplicity of computational hardware (e.g. multiple processor cores) in order to perform computation more quickly. Different parts of the computation are delegated to different processors that execute at the same time (in parallel).

Definition

Concurrency is a program-structuring technique in which there are multiple threads of control. The user sees their effects interleaved. Whether they actually execute at the same time or not is an implementation detail.

Concurrency vs. Parallelism

https://wiki.haskell.org/Parallelism_vs._Concurrency

Disclaimer: Not all programmers agree on the meaning!

Definition

A **parallel program** is one that uses a multiplicity of computational hardware (e.g. multiple processor cores) in order to perform computation more quickly. Different parts of the computation are delegated to different processors that execute at the same time (in parallel).

Definition

Concurrency is a program-structuring technique in which there are multiple threads of control. The user sees their effects interleaved. Whether they actually execute at the same time or not is an implementation detail.

CPU-bound tasks: multiprocessing/parallel processing

- Simultaneous. Like, really simultaneous!
- Multiple processes, multiple CPU cores
- Good for: matrix multiplications, searching, image processing, etc.
- Problems: computation parallelization, sending data between processes

CPU-bound tasks: multiprocessing/parallel processing

- Simultaneous. Like, really simultaneous!
- Multiple processes, multiple CPU cores
- Good for: matrix multiplications, searching, image processing, etc.
- Problems: computation parallelization, sending data between processes

CPU-bound tasks: multiprocessing/parallel processing

- Simultaneous. Like, really simultaneous!
- Multiple processes, multiple CPU cores
- Good for: matrix multiplications, searching, image processing, etc.
- Problems: computation parallelization, sending data between processes

CPU-bound tasks: multiprocessing/parallel processing

- Simultaneous. Like, really simultaneous!
- Multiple processes, multiple CPU cores
- Good for: matrix multiplications, searching, image processing, etc.
- Problems: computation parallelization, sending data between processes

I/O-bound tasks: threading, asynchronicity

- Only one process
- Not *really* simultaneous: but can sometimes fake it quite well
- Good for: sending requests (to DB, to APIs), waiting for user input, etc.
- Problems:

I/O-bound tasks: threading, asynchronicity

- Only one process
- Not *really* simultaneous: but can sometimes fake it quite well
- Good for: sending requests (to DB, to APIs), waiting for user input, etc.
- Problems:

I/O-bound tasks: threading, asynchronicity

- Only one process
- Not *really* simultaneous: but can sometimes fake it quite well
- Good for: sending requests (to DB, to APIs), waiting for user input, etc.
- Problems:

I/O-bound tasks: threading, asynchronicity

- Only one process
- Not *really* simultaneous: but can sometimes fake it quite well
- Good for: sending requests (to DB, to APIs), waiting for user input, etc.
- Problems:

Threading

- A **thread** of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler, which is typically a part of the operating system.
- Pre-emptive multitasking: the system (scheduler) interrupts threads at arbitrary moments, switches to another thread, and later resumes the stopped tasks
- No task cooperation is necessary
- Caveats: Race conditions! – random bugs. Thread-safety is important!

Threading

- A **thread** of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler, which is typically a part of the operating system.
- Pre-emptive multitasking: the system (scheduler) interrupts threads at arbitrary moments, switches to another thread, and later resumes the stopped tasks
- No task cooperation is necessary
- Caveats: Race conditions! – random bugs. Thread-safety is important!

Threading

- A **thread** of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler, which is typically a part of the operating system.
- Pre-emptive multitasking: the system (scheduler) interrupts threads at arbitrary moments, switches to another thread, and later resumes the stopped tasks
- No task cooperation is necessary
- Caveats: Race conditions! – random bugs. Thread-safety is important!

Threading

- A **thread** of execution is the smallest sequence of programmed instructions that can be managed independently by a scheduler, which is typically a part of the operating system.
- Pre-emptive multitasking: the system (scheduler) interrupts threads at arbitrary moments, switches to another thread, and later resumes the stopped tasks
- No task cooperation is necessary
- Caveats: Race conditions! – random bugs. Thread-safety is important!

Asynchronicity

- Event loop – definition
- Cooperative multitasking: tasks have to be programmed to yield control when they don't need system resources \Rightarrow easier resources sharing
- What if a task does not cooperate? – blocking operations

Asynchronicity

- Event loop – definition
- Cooperative multitasking: tasks have to be programmed to yield control when they don't need system resources \Rightarrow easier resources sharing
- What if a task does not cooperate? – blocking operations

Asynchronicity

- Event loop – definition
- Cooperative multitasking: tasks have to be programmed to yield control when they don't need system resources \Rightarrow easier resources sharing
- What if a task does not cooperate? – blocking operations

Built-ins

- `multiprocessing` – for CPU-bound tasks
- `threading` – ...?
- `asyncio` – from Python 3.5 on

Built-ins

- multiprocessing – for CPU-bound tasks
- threading – ...?
- asyncio – from Python 3.5 on

Built-ins

- multiprocessing – for CPU-bound tasks
- threading – ...?
- asyncio – from Python 3.5 on

Examples!

Global Interpreter Lock

- The mechanism used by the CPython interpreter to assure that only one thread executes Python bytecode at a time
- But why?!
- Makes the object model implicitly safe against concurrent access
- Reference counting – instead of garbage collection; (more about it in the next slide)

Global Interpreter Lock

- The mechanism used by the CPython interpreter to assure that only one thread executes Python bytecode at a time
- But why?!
- Makes the object model implicitly safe against concurrent access
- Reference counting – instead of garbage collection; (more about it in the next slide)

Global Interpreter Lock

- The mechanism used by the CPython interpreter to assure that only one thread executes Python bytecode at a time
- But why?!
- Makes the object model implicitly safe against concurrent access
- Reference counting – instead of garbage collection; (more about it in the next slide)

Global Interpreter Lock

- The mechanism used by the CPython interpreter to assure that only one thread executes Python bytecode at a time
- But why?!
- Makes the object model implicitly safe against concurrent access
- Reference counting – instead of garbage collection; (more about it in the next slide)

Global Interpreter Lock – continued

Reference counting: any reference to an object modifies it (or at least its refcount)

```
> import sys
> a = []
> sys.getrefcount(a)
2
> b = a
> sys.getrefcount(a)
3
```

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

- The reference count needs protection against race conditions!
- Otherwise: memory leaks (never released) or incorrectly released memory, while a reference to the object still exist
- **A solution?** Add locks to all the objects that are shared between threads
- *Consequences:* decreased performance, and deadlocks!
- **A better solution?** A single lock: execution of any Python code requires acquiring the lock on the interpreter
- *Consequences:* easier, thread-safe, but any CPU-bound program becomes, effectively, single-threaded.
- Back then, when Python was a young language, it made it easy to add C extensions (didn't have to be thread-safe) – this helped make Python more popular

Global Interpreter Lock – continued

The GIL is always released when doing I/O.

[O]nly the thread that has acquired the GIL may operate on Python objects or call Python/C API functions. In order to emulate concurrency of execution, the interpreter regularly tries to switch threads (...). The lock is also released around potentially blocking I/O operations like reading or writing a file, so that other Python threads can run in the meantime.

Global Interpreter Lock – continued

How to deal with GIL?

- multithreading → asyncio
- multithreading → multiprocessing (+ asyncio)
- CPU-intensive functions → Cython (no GIL!)
- Python → a faster language...?

Pycharm: concurrency diagrams (reveals locking issues, but omits GIL)

Global Interpreter Lock – continued

How to deal with GIL?

- multithreading → asyncio
- multithreading → multiprocessing (+ asyncio)
- CPU-intensive functions → Cython (no GIL!)
- Python → a faster language...?

Pycharm: concurrency diagrams (reveals locking issues, but omits GIL)

Global Interpreter Lock – continued

How to deal with GIL?

- multithreading → asyncio
- multithreading → multiprocessing (+ asyncio)
- CPU-intensive functions → Cython (no GIL!)
- Python → a faster language...?

Pycharm: concurrency diagrams (reveals locking issues, but omits GIL)

Global Interpreter Lock – continued

How to deal with GIL?

- multithreading → asyncio
- multithreading → multiprocessing (+ asyncio)
- CPU-intensive functions → Cython (no GIL!)
- Python → a faster language...?

Pycharm: concurrency diagrams (reveals locking issues, but omits GIL)

Concurrency in Python – bibliography & further reading

- Jim Anderson, "*Speed Up Your Python Program With Concurrency*",
<https://realpython.com/python-concurrency/>
- David Beazley, "*An Introduction to Python Concurrency*",
presented at USENIX Technical Conference San Diego, June, 2009. Slides available at
<https://speakerd.s3.amazonaws.com/presentations/3770713233254908b259542c4361e976/Concurrent.pdf>

GIL – bibliography & further reading

- Abhinav Ajitsaria, *"What is the Python Global Interpreter Lock (GIL)?"* <https://realpython.com/python-gil/>
- Python Wiki, *"Global Interpreter Lock"*,
<https://wiki.python.org/moin/GlobalInterpreterLock>
- *"Thread State and the Global Interpreter Lock"*,
<https://docs.python.org/3/c-api/init.html#thread-state-and-the-global-interpreter-lock>
- Christoph Heer, *"Is it me, or the GIL?"*, presented at EuroPython 2019 in Basel, Switzerland, July, 2019. Slides available at
<https://ep2019.europython.eu/media/conference/slides/Lj9n5pc-is-it-me-or-the-gil.pdf>